TASK 1 Self-Check:

Pick one part of your design that you are not sure about. Show us how you checked it. A few examples by hand, a small measurement, a quick probe - anything is fine. We want to see that you check your own work, not just ship it.

I was unsure whether LLM-generated test scripts (`test.sh`) would catch partial or fake fixes, or if they would falsely pass incomplete solutions.

Evaluated my 2 generated tasks
Hand-crafted 2 fake solutions for each task

1. Hardcoded expected output.
2. Partial fix that solves the basic path but breaks edge cases.

Executed the 2 fake solutions against their corresponding `test.sh` scripts.

Observation:
Dynamic inputs in `test.sh` prevented static mocking. For the partial fixes the generated tests only checked the happy path described in `instruction.md` and missed boundary conditions.

Action:
Based on this check, I updated the test-generator prompt to mandate at least two negative/edge-case assertion tests alongside the primary test path, eliminating false-pass vulnerabilities.

Additionally, My validation loop can be stronger for task generation, it is right now checking for the existence, but can also follow the stricter validation approaches.